

# Notebook

April 20, 2019




Now make a scatterplot using midpoints as the x variable and fraction\_outcome as the y variable. Draw a dashed line from [0,0] to [1,1] to mark the line  $y=x$ .

```
In [53]: %matplotlib inline
import matplotlib.pyplot as plt

plt.scatter(midpoints, fraction_outcome.values)
plt.title('Predicted vs Actual Plot')
plt.xlabel('Midpoints')
plt.ylabel('Actual win prob')
plt.plot([0,1], [0,1], '--')
```

```
Out[53]: [<matplotlib.lines.Line2D at 0x7f3a61d09c18>]
```





### 0.0.1 Question 5: adding error bars

If you did things correctly, it should look like fivethirtyeight has done “pretty” well with their forecasts: the actual fraction of wins tracks closely with the predicted number. But how do we decide what’s “good enough”? Consider this example: I correctly predict that a coin is fair (e.g. that it has a 50% chance of heads, 50% chance of tails). But if I flip it 100 times, I can be pretty sure it won’t come up heads exactly 50 times. The fact that it didn’t come up heads exactly 50 times doesn’t make my prediction incorrect.

To assess how reasonable the predictions are, I need to quantify the uncertainty in my estimate. It’s reasonable to assume that within each bin,  $k$ , the observed number of wins,  $Y_k \sim \text{Bin}(n_k, p_k)$ , where  $n_k$  is the number of elections and  $p_k$  is the predicted win probability in bin  $k$ .

Classical results tell us that the observed fraction of wins in bin  $k$ ,  $\hat{p} = \frac{Y_k}{n_k}$  has variance  $\text{Var}(\hat{p}_k) = \frac{p_k(1-p_k)}{n_k} \approx \frac{\hat{p}_k(1-\hat{p}_k)}{n_k}$ . The standard deviation of the Binomial proportion then is  $\hat{\sigma}_k \approx \sqrt{\frac{\hat{p}_k(1-\hat{p}_k)}{n_k}}$ .

If we use the [normal approximation to generate a confidence interval](#), then the 95% interval has the form  $\hat{p}_k \pm 1.96\hat{\sigma}_k$ .

Create a new “aggregated” dataframe. This time, group `election_sub` by the bin and compute both the average of the `probwin_outcome` (mean) and the number of observations in each bin (count) using the `agg` function. Call this new data frame, `election_agg`.

```
In [55]: election_agg = election_sub.groupby('bin')['probwin_outcome'].agg(['mean', 'count'])
         election_agg
```

```
Out[55]:
```

	mean	count
bin		
(0.0, 0.0526]	0.001715	583
... 0mitting 14 lines ...		
(0.842, 0.895]	0.976744	43
(0.895, 0.947]	0.937500	32
(0.947, 1.0]	0.998478	657



Use the mean and count columns of election\_agg to create a new column of election\_agg titled err, which stores  $1.96 \times \hat{\sigma}_k$  in each bin  $k$ .

```
In [56]: #election_agg['count']
         var = (election_agg['mean']*(1-election_agg['mean']))/election_agg['count']
         error = (var**0.5)*1.96
         election_agg['err'] = error
         election_agg
```

```
Out[56]:
```

	mean	count	err
bin			
(0.0, 0.0526]	0.001715	583	0.003359
... 0mitting 14 lines ...			
(0.842, 0.895]	0.976744	43	0.045048
(0.895, 0.947]	0.937500	32	0.083870
(0.947, 1.0]	0.998478	657	0.002981





### 0.0.2 Question 7: understanding confidence intervals

Are the 95% confidence intervals generally larger or smaller for more confident predictions (e.g. the predictions closer to 0 or 1). What are the factors that determine the length of the confidence intervals?

**SOLUTION HERE** Higher confidence indicates a larger interval, since you are more confident that the true value will fall between a larger range. Sample size is one big factor that can determine the length of the confidence intervals as well as overall standard error, i.e. the variance. A smaller sample size would give you a larger interval.



```
In [60]: # Input: a pandas dataframe with a numeric column named `probwin`  
# Output: a pandas dataframe with the same columns, with an additional column named `absdiff`  
def abs_diff(x):  
    x['absdiff'] = max(x['probwin']) - min(x['probwin'])  
    return x
```



We can use this function to compute the difference between the maximum and minimum predicted with probabilities for every candidate. To do so, group `election_sub` by candidate and apply the function `abs_diff`. Find the index of the largest difference in `diff_dataframe` and store it in `max_idx`. Do this using `np.nanargmax` function. This function finds the *index* of the largest value, ignoring any missing values (nans).

```
In [66]: diff_dataframe = election_sub.groupby('candidate').apply(abs_diff)
diff_dataframe
max_idx = np.nanargmax(diff_dataframe['absdiff'])
#diff_dataframe.iloc[max_idx]
```

```
Out[66]: year                2018
office                House
state                 KS
... Omitting 10 lines ...
bin                (0.842, 0.895]
absdiff                0.65428
Name: 1890, dtype: object
```



Did the candidate win or lose the election?

```
In [71]: election_sub[(election_sub.candidate == 'Sharice Davids')]
```

```
Out[71]:
```

	year	office	state	district	special	election_date	forecast_date	\
1890	2018	House	KS	3.0	False	2018-11-06	2018-11-06	
252855	2018	House	KS	3.0	False	2018-11-06	2018-08-11	
... Omitting 4 lines ...								
		actual_voteshare	probwin	probwin_outcome			bin	
1890		NaN	0.84994		1	(0.842, 0.895]		
252855		NaN	0.19566		1	(0.158, 0.211]		






Now create a lineplot with forecast date on the x-axis and the predicted win probability on the y-axis.

```
In [76]: predicted_probs.plot.line(x='forecast_date', y='probwin')
plt.title('Predicted win probability for Sharice Davids over time')
plt.xlabel('Forecast date')
plt.ylabel('Predicted win probability')
```

```
Out[76]: Text(0, 0.5, 'Predicted win probability')
```






### 0.0.3 Question 10: prediction histograms

Make a histogram showing the predicted win probabilities on the morning of the election. Again, restrict yourself to only the classic predictions.

```
In [77]: electionplot = election_data[(election_data.forecast_type=="classic") & (election_data.forecast_type=="classic")]
electionplot['probwin'].hist()
plt.title('Histogram of predicted win probabilities on election day')
plt.xlabel('Predicted win probabilities')
plt.ylabel('Frequency')
```

```
Out[77]: Text(0, 0.5, 'Frequency')
```





Are most house elections easy to forecast or hard to forecast?

**SOLUTION HERE** House elections are easy to forecast because there is not a lot of randomness associated with predicting them, as it can be said that most incumbents will most likely vote to keep their seat. Given other situations, predicted house elections does not possess as much randomness as would be thought, as there is a better sense of what affects outcomes.



Create a pandas dataframe from the csv and print the first 10 rows.

```
In [80]: baseball_data = pd.read_csv("mlb_games.csv")
        baseball_data.head(10)
```

```
Out[80]:
```

	season	date	team1	team2	dh	prob1	prob1_outcome	\
0	2018	2018-10-28	Dodgers	Red Sox	0	0.483877	0.0	
1	2018	2018-10-27	Dodgers	Red Sox	0	0.508342	0.0	
... Ommitting 16 lines ...								
7	0.396971		1.0					
8	0.431984		1.0					
9	0.399572		0.0					






In this dataframe prob1 is the predicted win probability for team1. Make a histogram of prob1. Set the limits of the x-axis to [0, 1]

```
In [79]: baseball_data['prob1'].hist(range = [0,1])  
        plt.title('Histogram of predicted win probability for team 1')  
        plt.xlabel('Predicted win probability')  
        plt.ylabel('Frequency')
```

```
Out[79]: Text(0, 0.5, 'Frequency')
```





Find the most “surprising” baseball game outcome. To do so, select all of the entries for which `prob1_outcome` is 1 (i.e. team1 won the game), and then look for the index of the row containing the smallest value of `prob1`. This will correspond to the game that was most surprising according to `fivethirtyeights` predictions. Find and print the row corresponding to this most surprising outcome.

```
Out[82]:
```

	season	date	team1	team2	dh	prob1	prob1_outcome	\
521	2018	2018-08-25	Royals	Indians	0	0.287558		1.0
	prob2	prob2_outcome						
521	0.712442	0.0						



### 0.0.5 Question 13

Are the outcomes of baseball games generally easier or harder to predict than the outcomes of political elections? In a few sentences, comment on why this might be the case. What data is available for these predictions? What factors affect the outcomes of elections and baseball games? What makes an event like an election or a baseball game “random”?

**SOLUTION HERE** Baseball games outcomes are generally harder to predict than political elections. There is more randomness involved with baseball games, as in, not everything in baseball games are clearly predicted and defined. For example, there are many factors in a baseball game that can affect the outcome, i.e. the players playing in that specific game, stamina, etc. Therefore, in order to better predict baseball games, you would need more data than just probability of that team winning. These are also the events that make an election or a baseball game random, because we cannot thoroughly predict and account for these factors.



Create an analogous plot for empirical error bars with `bootstrap_election_agg`. Also draw a horizontal lines at 0 and 1.

**SOLUTION HERE**





Compare the two error bar plots and explain.  
**SOLUTION HERE**